Operators Monitoring and Observability Best Practices in Operator SDK

Shirly Radco
Principal Software Engineer
Red Hat

João Vilaça
Associate Software Engineer
Red Hat

February 2023

What Will We Talk About Today

Operators Observability

- When to start
- Metrics - Levels of Maturity
- Why monitor
- What to Monitor
- Best Practices in Operator SDK
- Code Examples in Memcached Operator

When should I start learning about monitoring

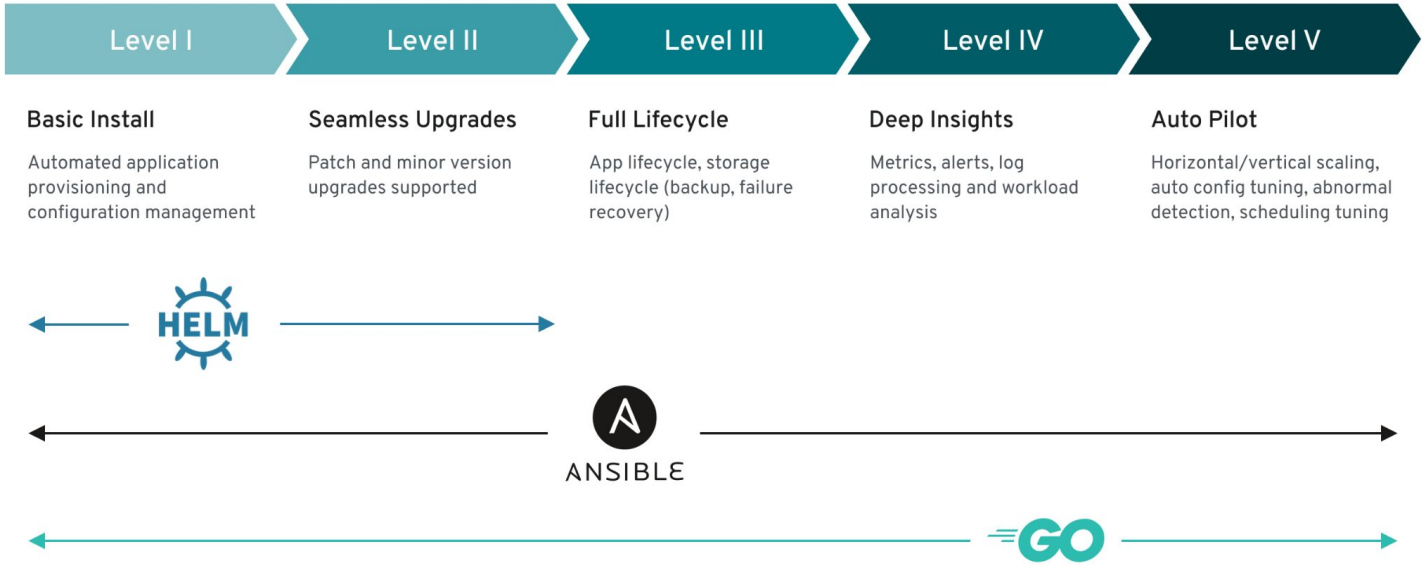

Source: https://sdk.operatorframework.io/docs/overview/operator-capabilities/

When should I start learning about monitoring

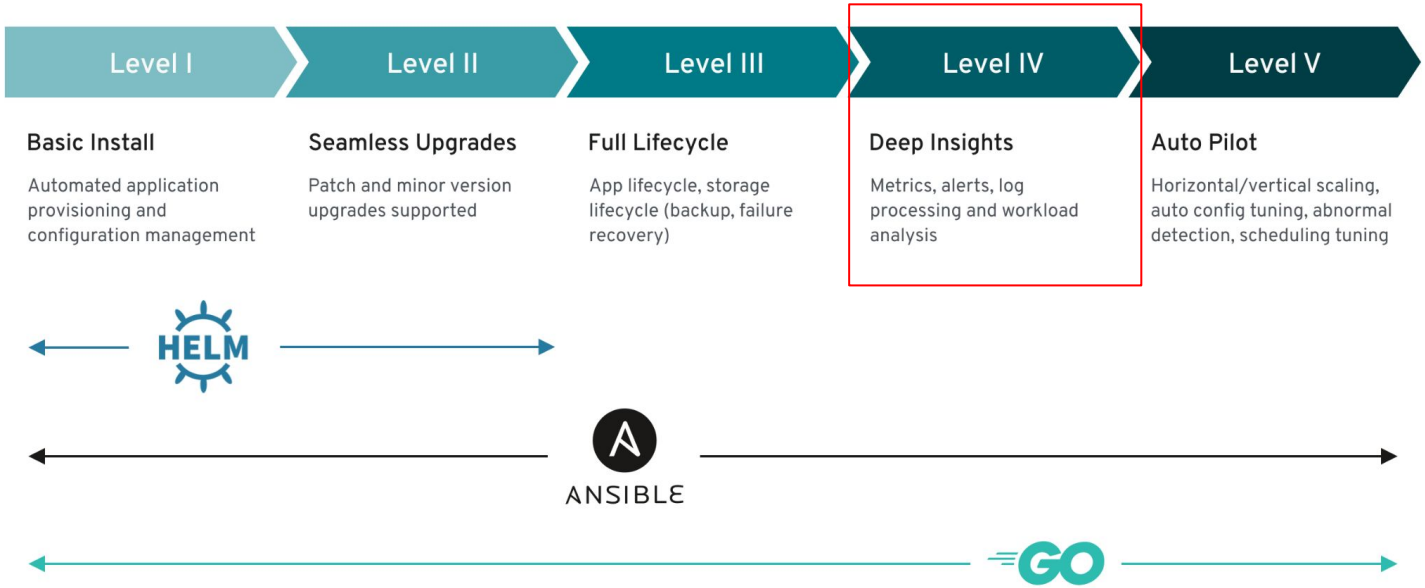

Source: https://sdk.operatorframework.io/docs/overview/operator-capabilities/

When should I start learning about monitoring

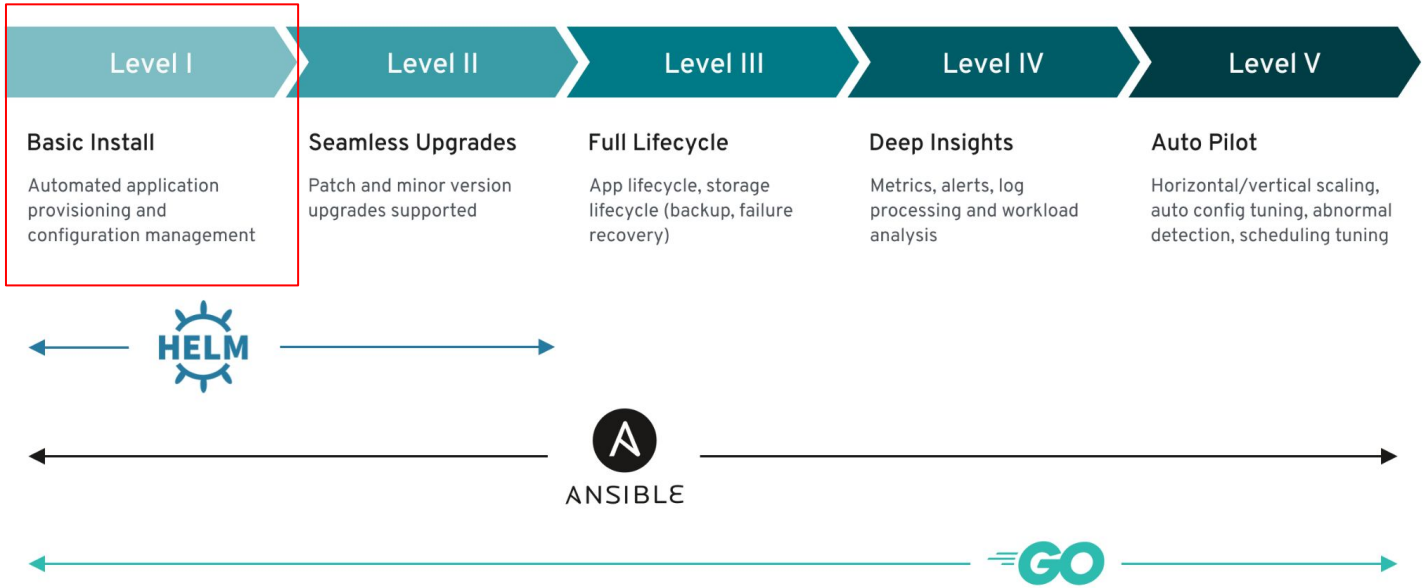

Source: https://sdk.operatorframework.io/docs/overview/operator-capabilities/

Operator Metrics - Levels of Maturity

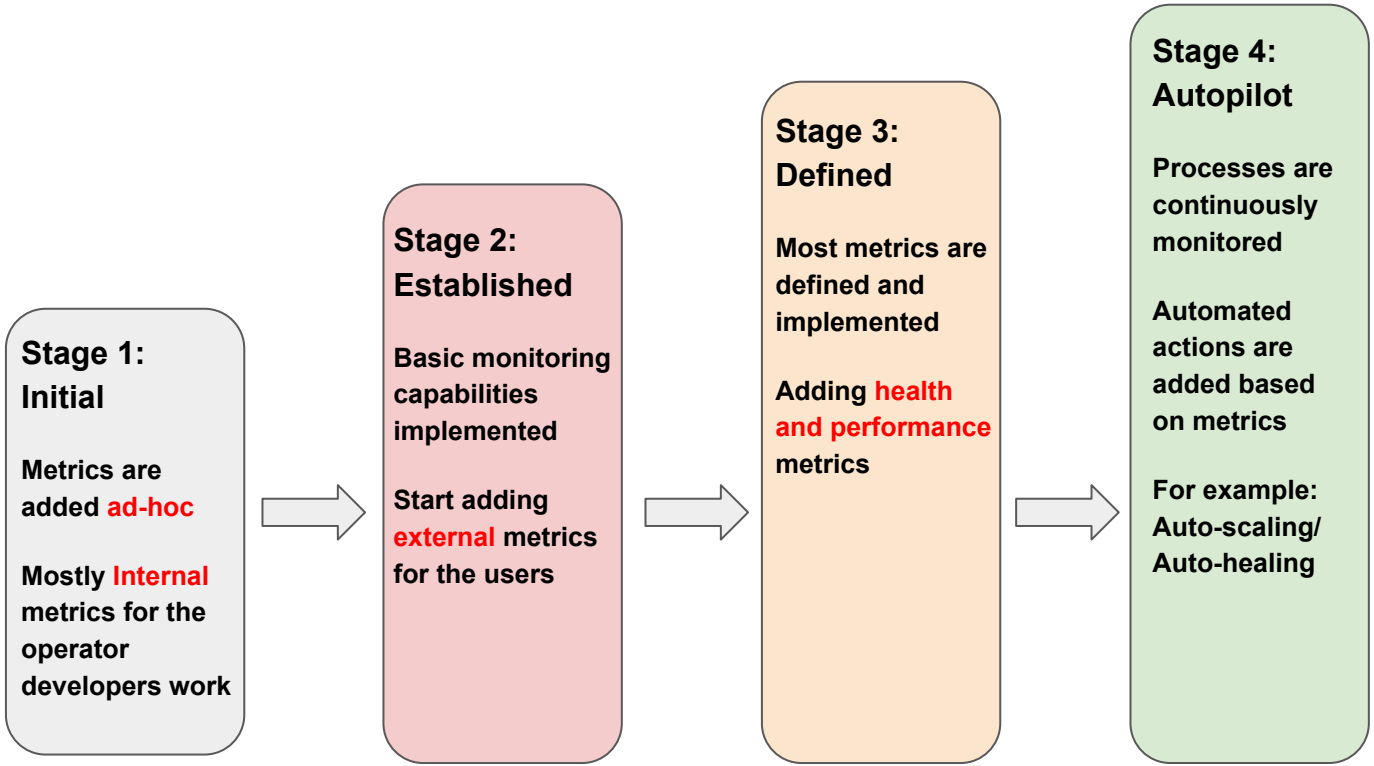

What Should Operator Developers Monitor?

The Operator and Custom Resources **Performance and Health**

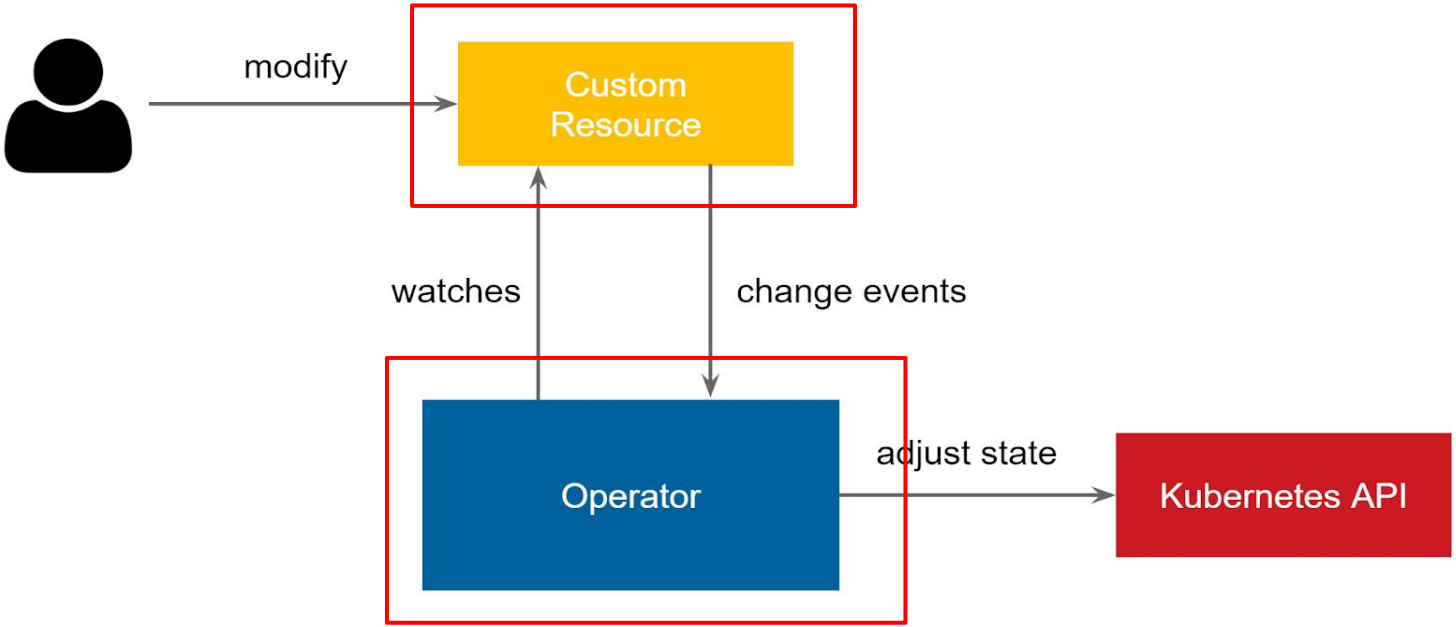

Source:https://cloud.redhat.com/blog/build-your-kubernetes-operator-with-the-right-tool

Why Add Operator Observability?

Performance and Health

Detect issues early, reduce downtime, detect regressions

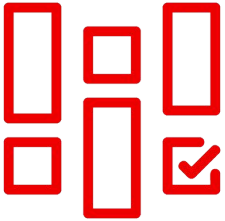

Planning and Billing

Improve resources planning and Profitability

Alerts

Enable **Proactive actions** and **Educate the user**

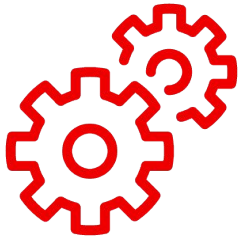

Autopilot

Auto-scaling, Auto-healing, Auto-tuning, Abnormality detection

Operator Observability in Operator SDK

- **Tools to build, test and package Operators**
- **Leading practices and code patterns to help prevent reinventing the wheel.**

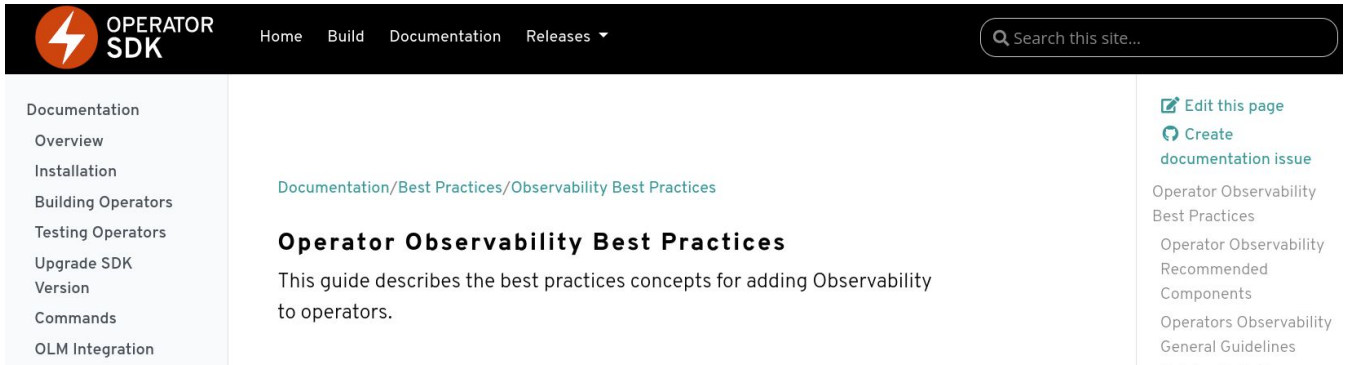

https://sdk.operatorframework.io/docs/best-practices/observability-best-practices/

Observability Best Practices - Objectives

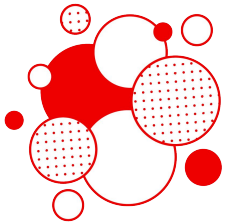

The Power of Community

Learn from others mistakes and **avoid known pitfalls**

Shorten on-boarding and coding time

Code examples and documentation for Observability in operators

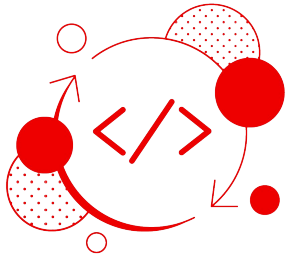

Reusable Code

Saves Time, Lower Cost, Reduced Development Risks, Prevents Code Bloat

Operator Metrics - Naming Conventions

Operator Metrics Name Prefix - Helps to search for all operator metrics

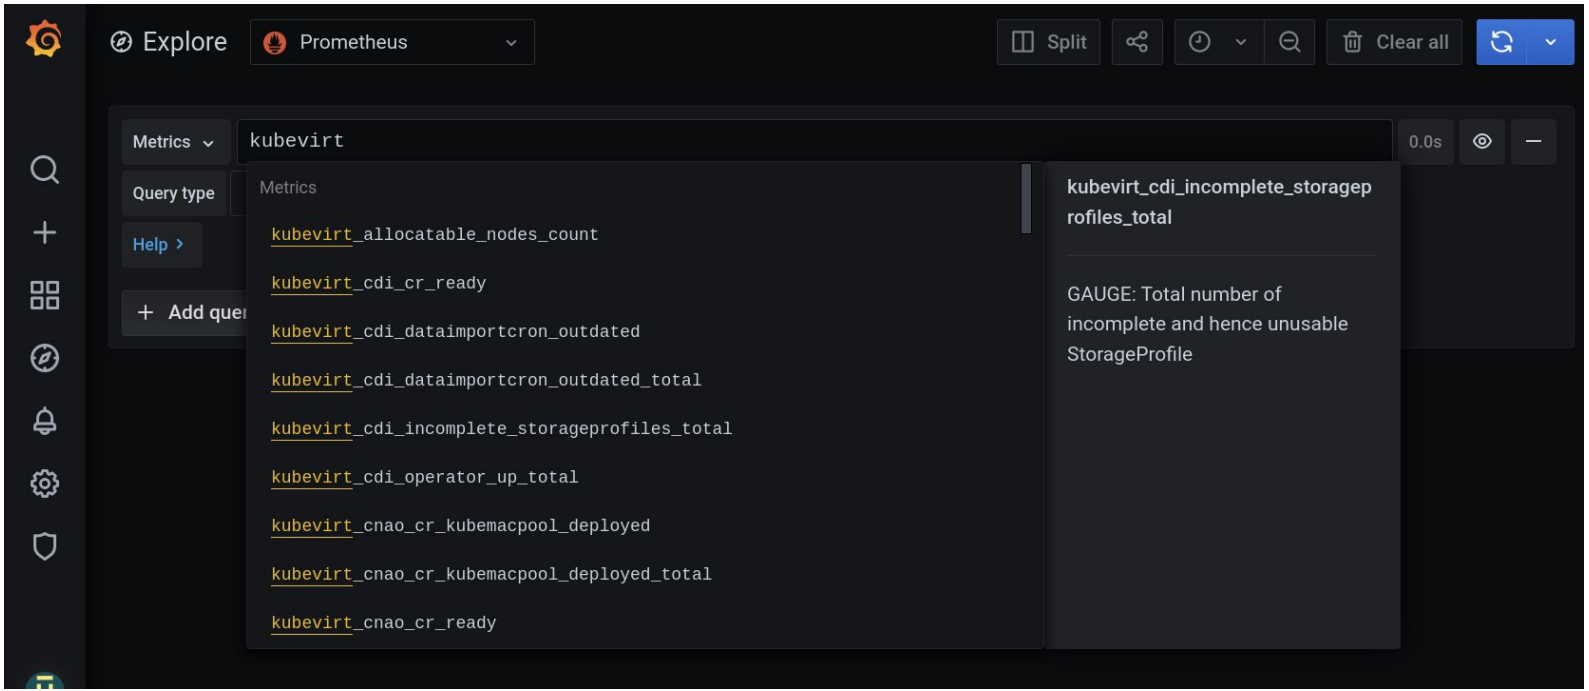

Operator Metrics - Naming Conventions

Operator Metrics Name Prefix - No prefix.

- Hard to understand who generates this (cAdvisor)
- Can’t search all cAdvisor metrics

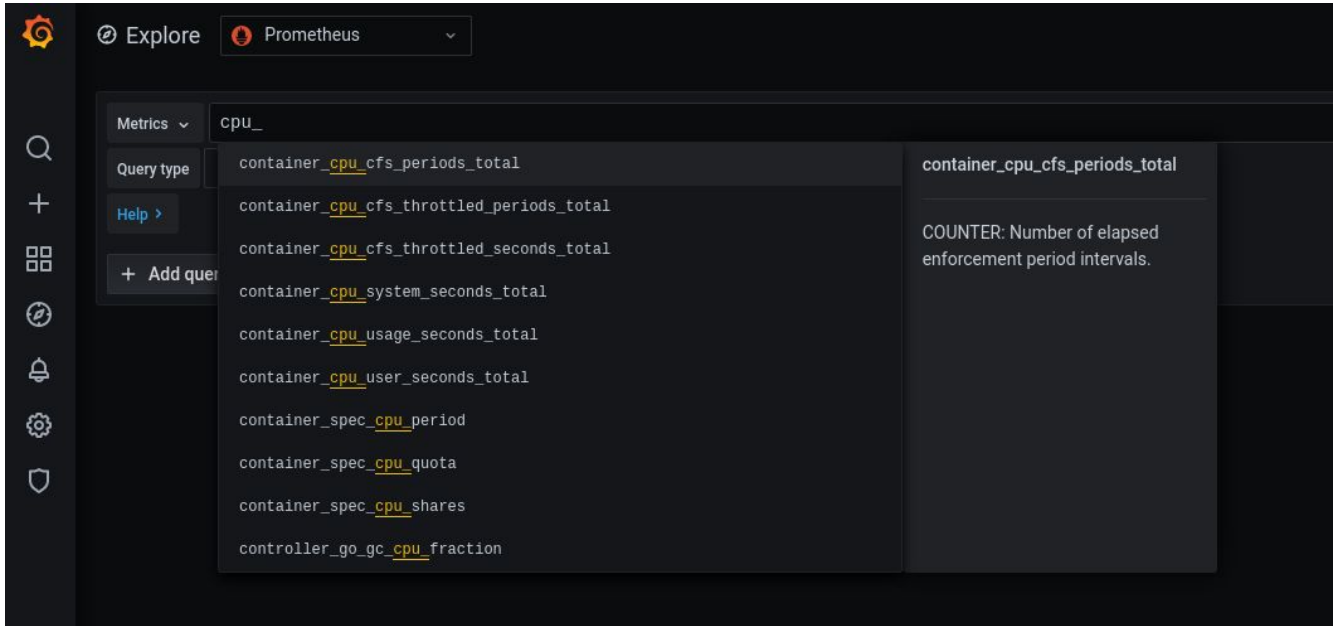

Operator Metrics - Help Text

Operator Metric Help Text - When adding a new metric, add a meaningful help text

- Helps users to **learn about the metric**
- Can be used for creating **auto generated documentation**

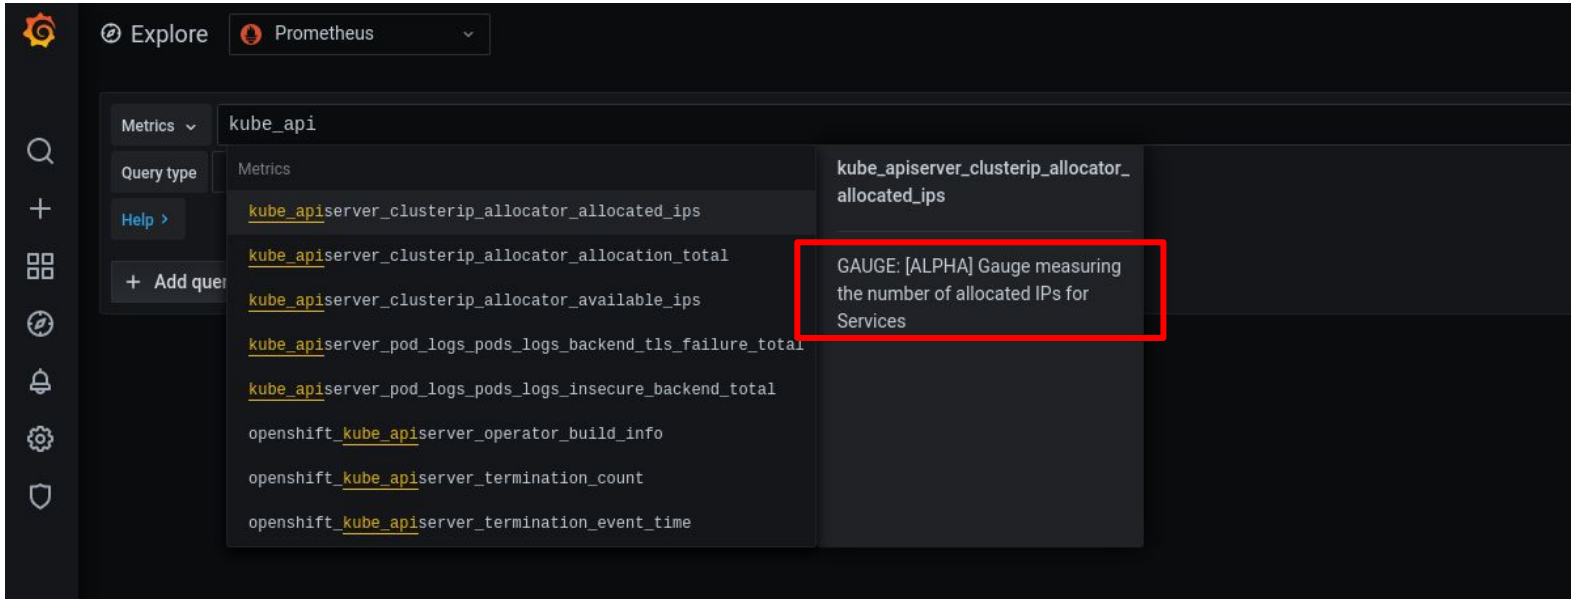

Operator Metrics - Base Units

Base units

Prometheus does not have any units hard coded. For better compatibility, base units should be used. The following lists some metrics families with their base unit. The list is not exhaustive.

Family	Base unit	Remark
Time	seconds	
Temperature	celsius	*celsius* is preferred over *kelvin* for practical reasons. *kelvin* is acceptable as a base unit in special cases like color temperature or where temperature has to be absolute.
Length	meters	
Bytes	bytes	
Bits	bytes	To avoid confusion combining different metrics, always use *bytes*, even where *bits* appear more common.
Percent	ratio	Values are 0–1 (rather than 0–100). `ratio` is only used as a suffix for names like `disk_usage_ratio`. The usual metric name follows the pattern `A_per_B`.
Voltage	volts	
Electric current	amperes	
Energy	joules	
Power		Prefer exporting a counter of joules, then `rate(joules[5m])` gives you power in Watts.
Mass	grams	*grams* is preferred over *kilograms* to avoid issues with the *kilo* prefix.

https://prometheus.io/docs/practices/naming/#base-units

Operator Metrics - Base Units

- The Prometheus approach is to **use base units** such as seconds, and then **floating point numbers** to store the value.
- Floating point **removes the concern of the magnitude of the number**

https://www.robustperception.io/who-wants-seconds/

Operator Metrics - Base Units

✓ **Operator Metrics Units** - Uses Prometheus base units “seconds”

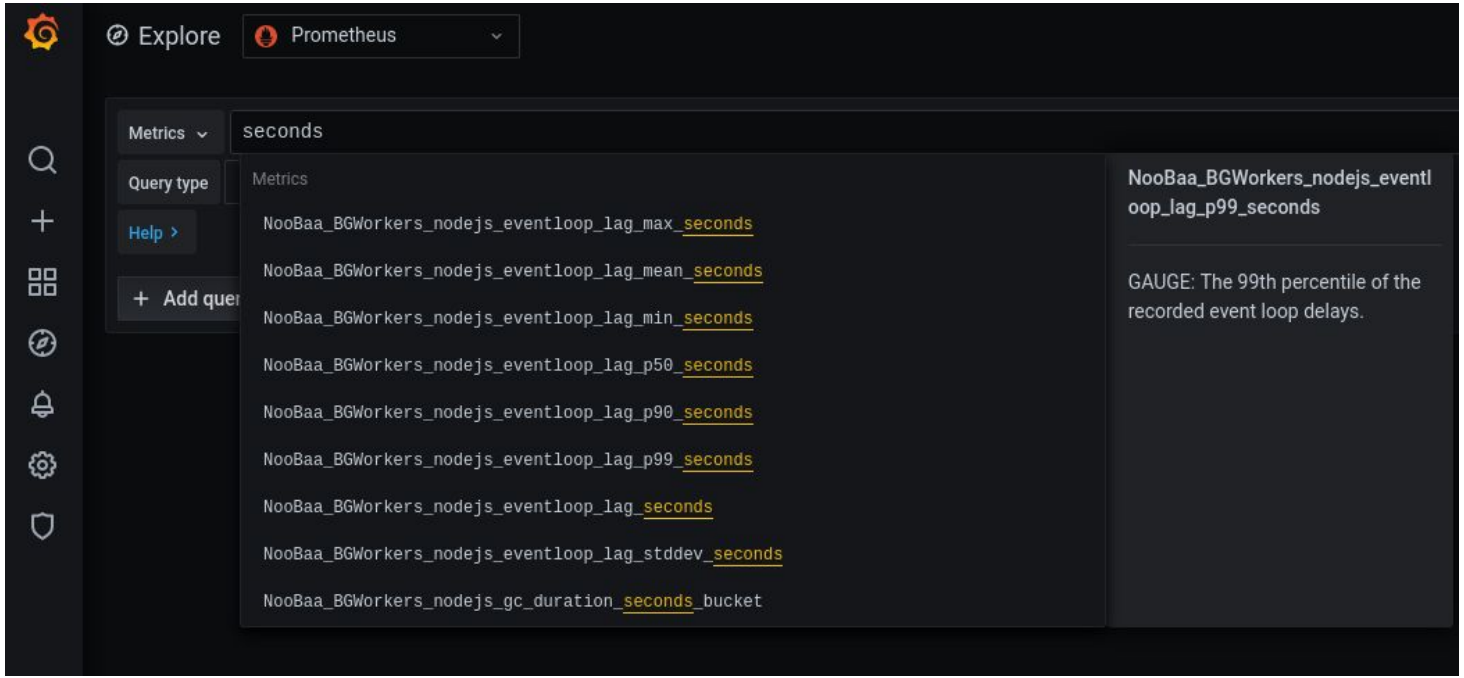

Operator Metrics - Base Units

Operator Metrics Units - Uses “milliseconds” instead of “seconds”

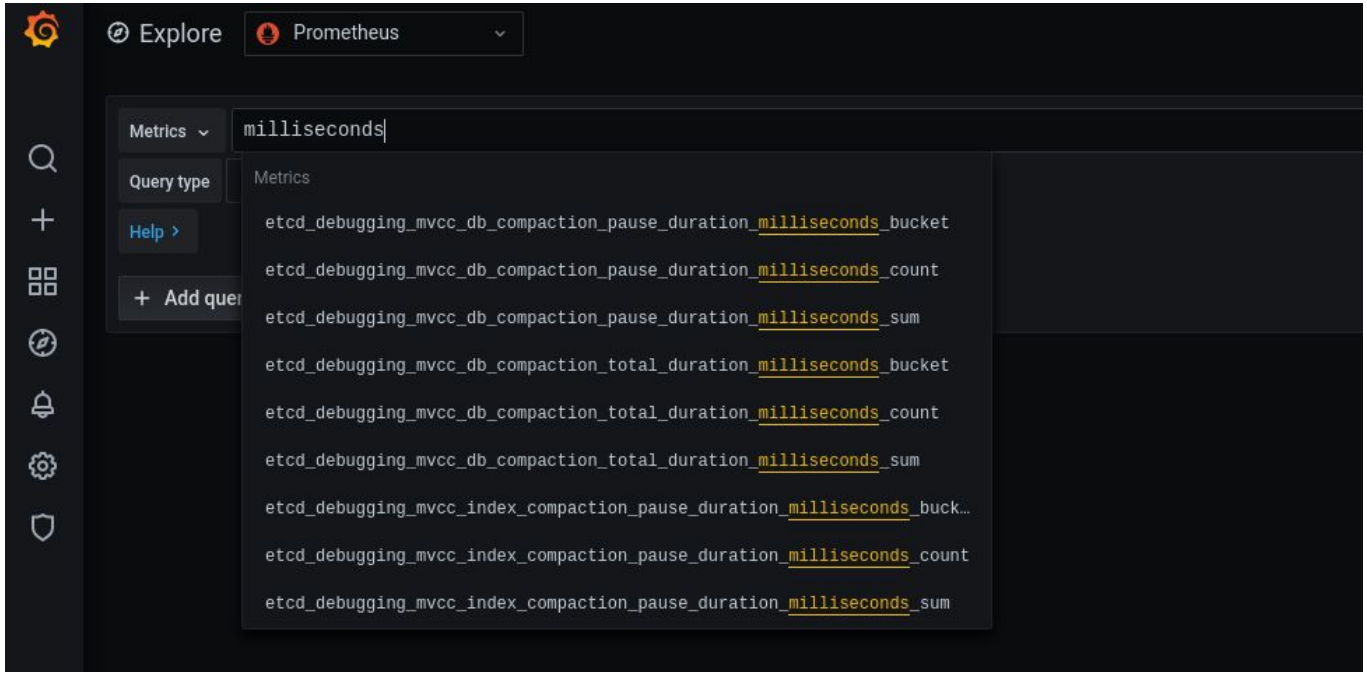

Operator Alerts - Example

```
// createOperatorDownAlertRule creates MemcachedOperatorDown alert rule
func createOperatorDownAlertRule() monitoringv1.Rule {
        return monitoringv1.Rule{
                Alert: operatorDownAlert,
                Expr:  intstr.FromString("memcached_operator_up_total == 0"),
                Annotations: map[string]string{
                        "description": "No running memcached-operator pods were detected in the last 5 min.",
                },
                For: "5m",
                Labels: map[string]string{
                        "severity":    "critical",
                        "runbook_url": runbookURLBasePath + "MemcachedOperatorDown.md",
                },
        }
}
```

Operator Alerts - Severity Label

Operator Alerts Severity - Only 3 valid Severity Label values : **Critical** / **Warning** / **Info**

Note: “none” severity - Only for the Watchdog

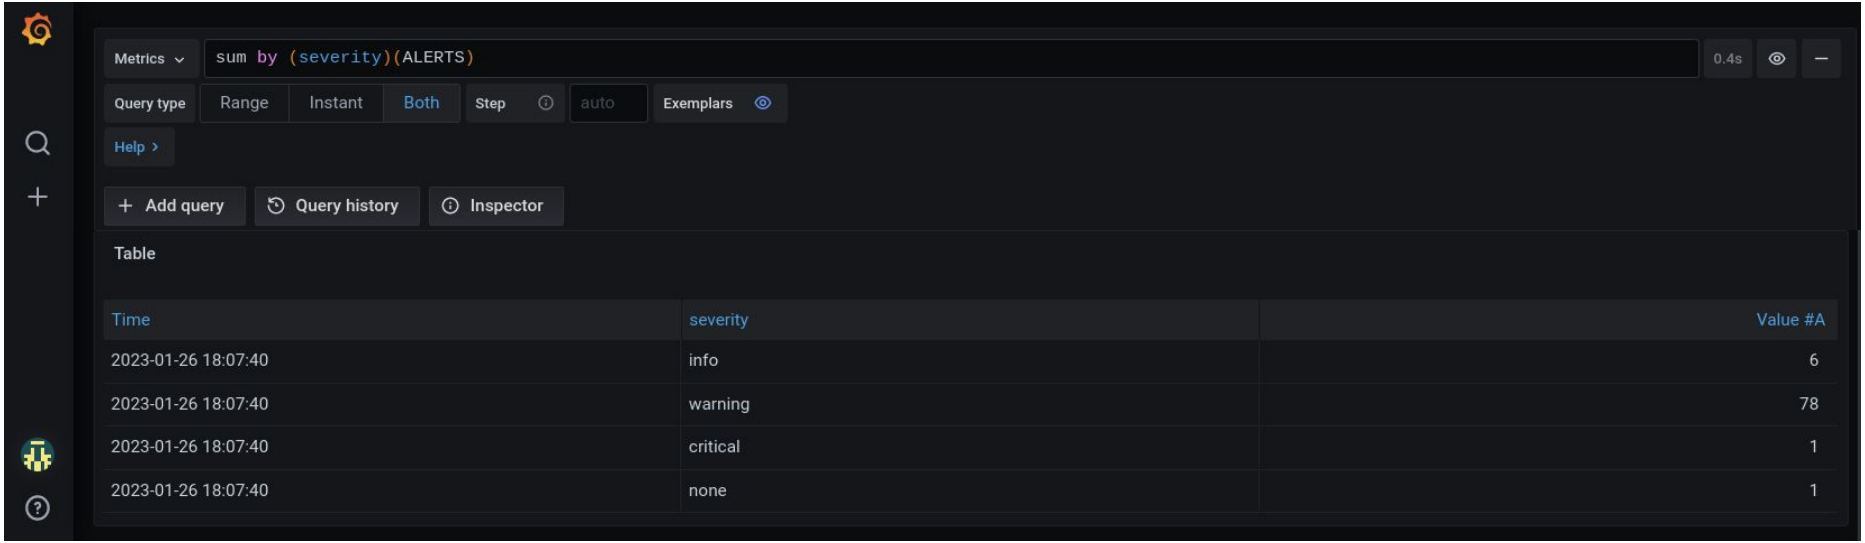

Operator Alerts - Severity Label

Operator Alerts Severity Label - “Major“ is not a valid value for Severity label

```
- alert: ScrapeProblem
  expr: up{kubernetes_namespace!~"openshift-.+",kubernetes_pod_name=~".+-kafka-[0-9]+"} == 0
  for: 3m
  labels:
    severity: major
  annotations:
    summary: 'Prometheus unable to scrape metrics from {{ $labels.kubernetes_pod_name }}/{{ $labels.instance }}'
    description: 'Prometheus was unable to scrape metrics from {{ $labels.kubernetes_pod_name }}/{{ $labels.instance }} for more than 3 minutes'
```

* Red Hat AMQ Streams

Operators Observability - Best Practices Status

Today:

- **Metrics** naming convention
- **Metrics documentation**
- **Alerts** required labels
- **Alert runbooks**
- Metrics and Alerts **tests**

Future:

- **Dashboards**
- **Logging**
- **Tracing**
- **Events**

Managing Code Complexity

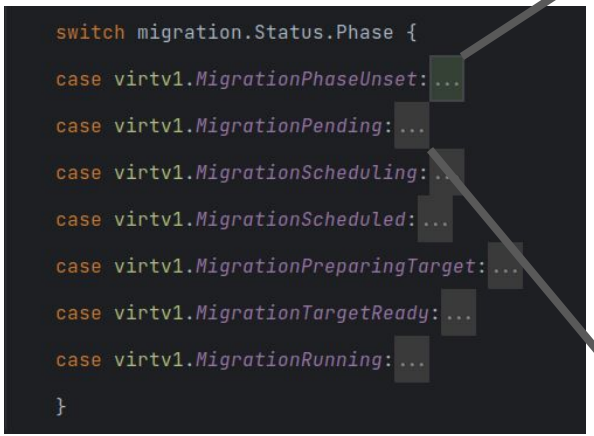

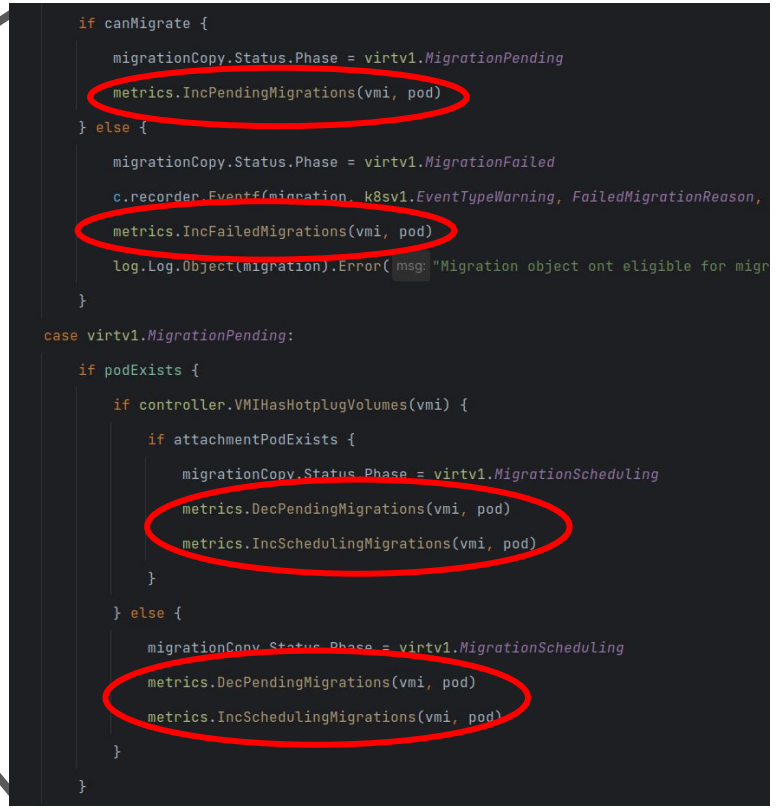

https://github.com/kubevirt/kubevirt ->
pkg/virt-controller/watch/migration.go

16.4 %

monitoring code intertwined in migration logic code

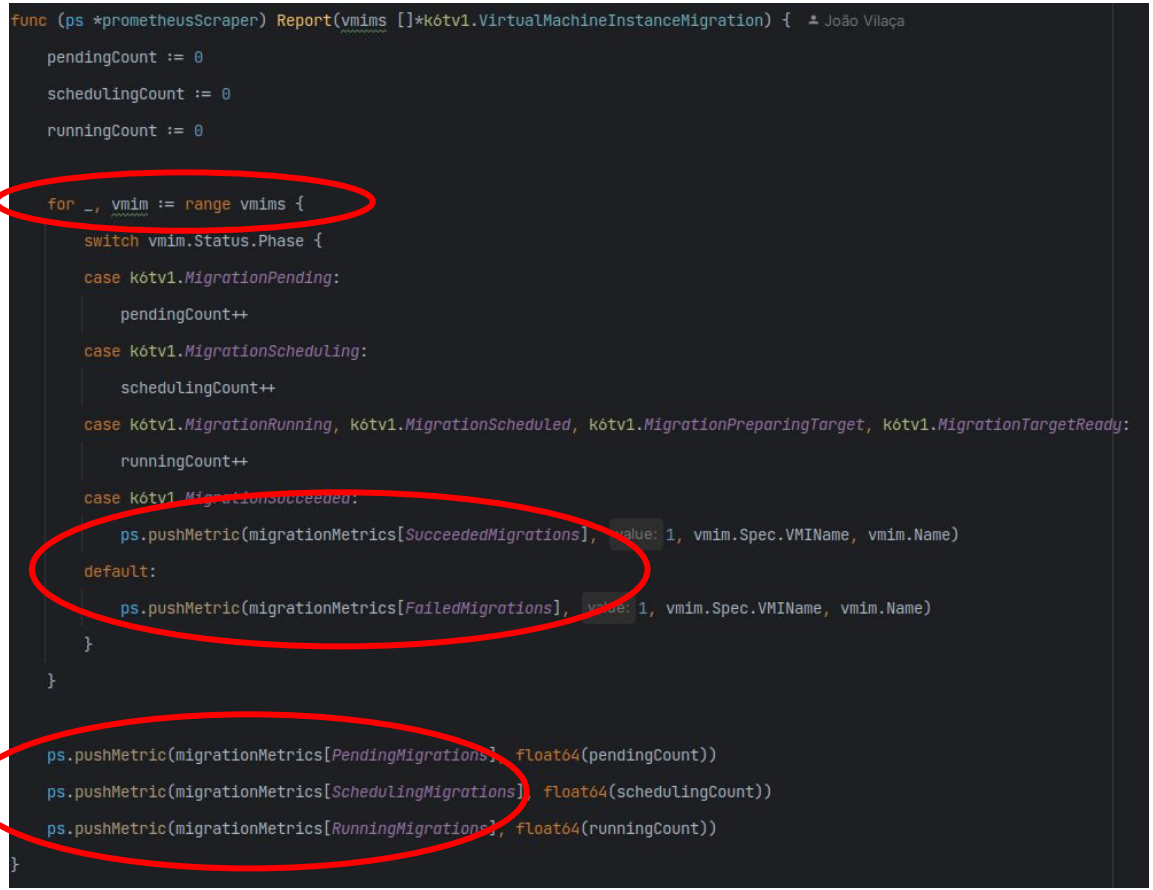

Usage of collector approach for VM migration metrics

Memcached Example Operator

An example operator of building a simple operator using tools and libraries provided by the Operator SDK.

https://sdk.operatorframework.io/docs/building-operators/golang/quickstart/

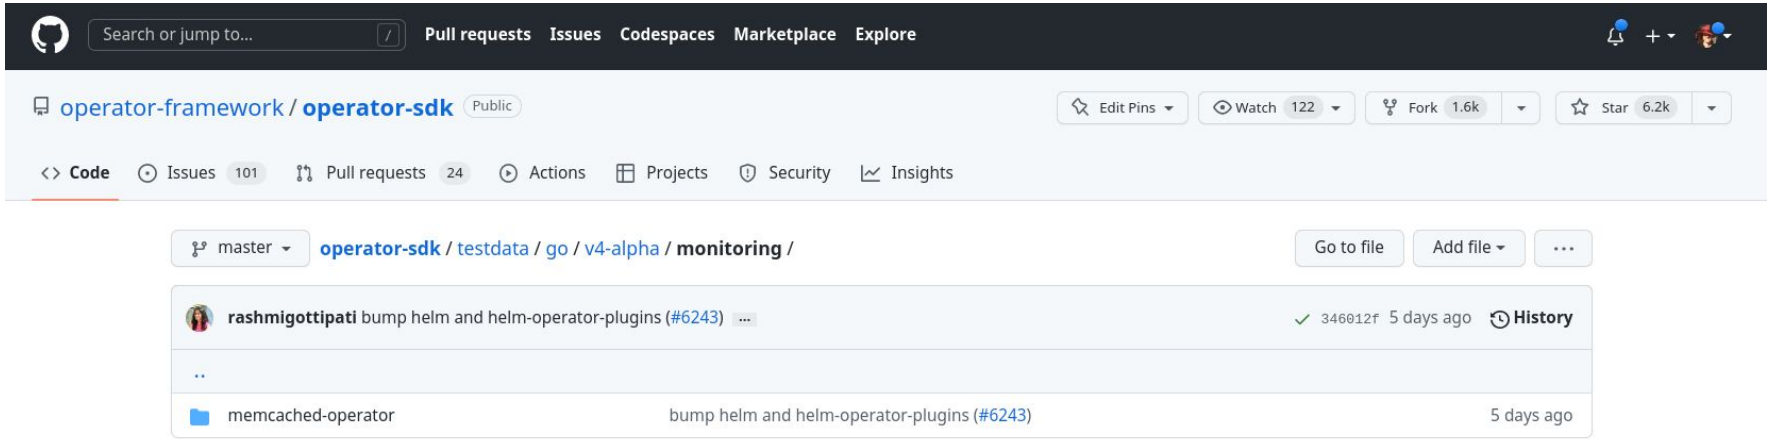

https://github.com/operator-framework/operator-sdk/tree/master/testdata/go/v4-alpha/monitoring/memcached-operator

Operator Observability Code

Clear

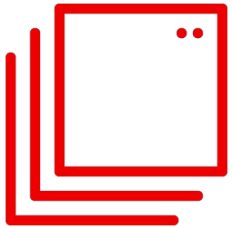

Modular

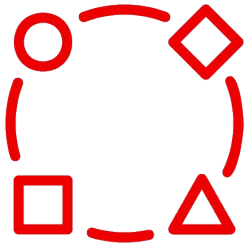

Reusable

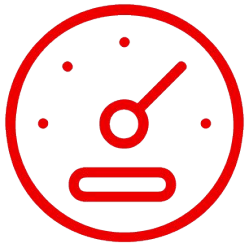

Performant

Operator Observability Code

Keep Monitoring Code Separate from Logic Code

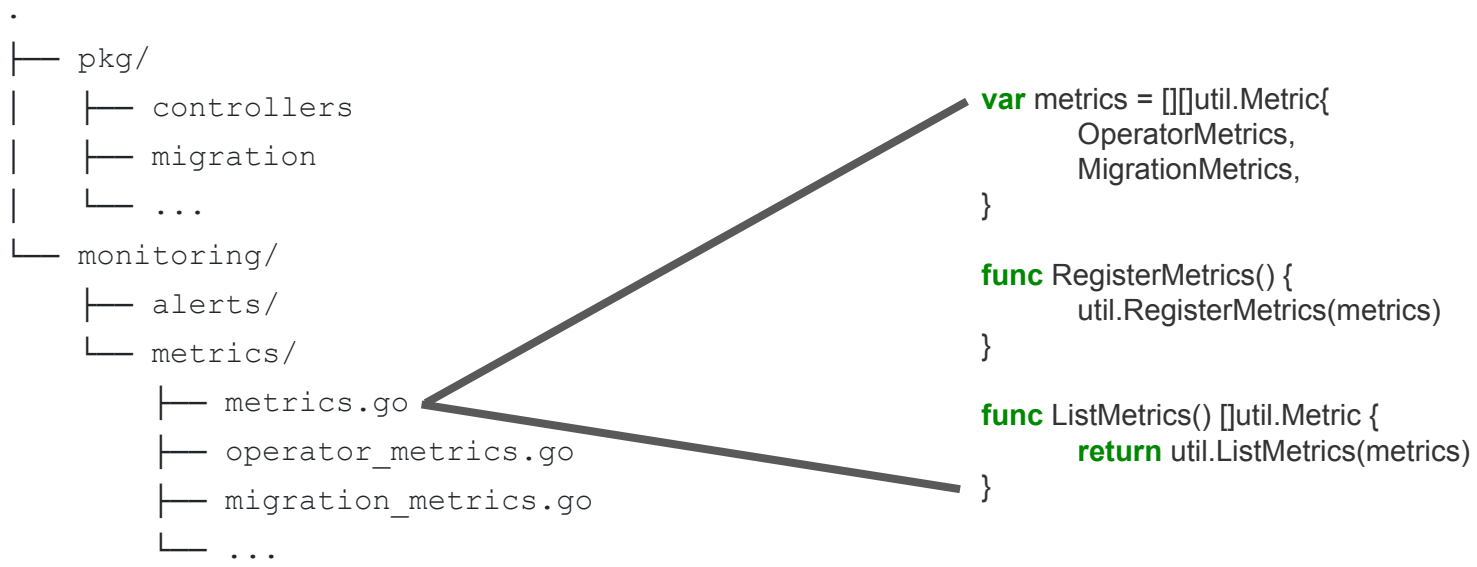

Operator Observability in Memcached Operator

Automate metric and alert code generation

Metrics name linter

Automated metrics documentation

Easy structure for unit and e2e testing

Thank you!

Shirly Radco
Principal Software Engineer
Red Hat

João Vilaça
Associate Software Engineer
Red Hat

February 2023